

Harsha Manjunatha

857-421-8166 | Boston, MA | harsha.manjunatha98@gmail.com | [LinkedIn](#) | [Portfolio](#) | [Github](#)

SUMMARY

Data Engineer with 5+ years of experience across supply chain, fintech and gaming industries, specializing in scalable ETL pipelines, cloud database management, and workflow orchestration. Proficient in **Python**, **SQL**, **Oracle SQL**, **Apache Airflow**, **AWS** and **GCP** with hands-on experience building, monitoring, and optimizing data infrastructure. Track record of improving pipeline efficiency, data availability, and query performance across transactional and analytical workloads.

EXPERIENCE

FreshDirect, *Bronx, New York*

Oct 2025 - Mar 2026

Supply Chain Strategy Data Analyst (Co-op)

- Recovered **\$400K+** in annual billback revenue by auditing vendor billing logic across all product categories, tracing systematic Oracle SQL calculation errors to a date-field mismatch and correcting the underlying pipeline.
- Engineered an **SQL**-based promotional calculations ETL pipeline replacing legacy Excel workflows, cutting query latency by **90%** and enabling real-time analytics across 14 product categories.
- Applied **K-Means clustering** (Python, scikit-learn) to segment 10,000+ SKUs into 4 performance tiers, driving targeted inventory and markdown strategies across 8 seasonal events.
- Integrated **Claude (Opus 4.6) LLM** via **MCP** server into a **Power BI** and SQL pipeline to automate billback and rebate forecast tracking across 14 product categories, delivering interactive dashboards with daily trends and YoY comparisons for annual vendor negotiations.
- Conducted A/B test analysis across 6 customer cohorts using **DiD regression** and **z-tests**, identifying a 1.4-2.5pp shift in express routing and **~1%** net order loss, to guide leadership's rollout decision.

IDFC First Bank, *Bangalore, India*

Jan 2022 - Aug 2024

Data Engineer

- Engineered real-time data pipelines using **Apache Kafka** and **Spark Streaming** to ingest JSON data from AWS MSK into HDFS, enhancing data availability by 25% and ensuring robust ETL processes with diligent data validation.
- Built monitoring dashboard using **SQL** and **Python** with automated validations and 30-minute email alerts (CloudWatch-integrated), enabling proactive pipeline issue detection and resolution.
- Collaborated on dynamic **Apache Airflow** DAG wrappers to facilitate incremental data fetching, staging, encryption, and seamless loading into Hive and **AWS S3**, improving overall pipeline efficiency by 30%.
- Partnered with the migration team to restructure and optimize 5+ data engineering projects during the transition from Kubernetes clusters to the **Cloudera** platform, boosting system performance by 30% and scalability by 40%.
- Automated complex SQL queries for daily airport lounge access reporting, securely distributing results via SFTP to stakeholders for timely, data-driven decision-making.

Pole To Win, *Bangalore, India*

Oct 2020 - Jan 2022

Data Analyst

- Enhanced real-time dashboards in **Tableau** to monitor key metrics (DAU, MAU, ARPU, retention rates), facilitating data-driven decision-making and improving player retention by 15%.
- Leveraged **SQL** and **Python** (Pandas, NumPy) to aggregate and transform 10M+ rows of game data from Amazon S3, ensuring data accuracy and performing exploratory analysis to identify actionable trends.
- Coordinated with cross-functional teams including product managers, marketing, and game designers to deliver interactive visualizations that drove a 20% increase in player engagement.

TECHNICAL SKILLS

Languages & Processing: Python (Pandas, NumPy, Matplotlib, scikit-learn), SQL, PySpark, SparkSQL

Databases & Warehouses: PostgreSQL, Oracle, MySQL, Hive, HDFS, Snowflake, BigQuery

Data Engineering: Apache Kafka, Apache NiFi, Apache Airflow, Apache Spark, Hadoop, Hive, DBT

Cloud Services: AWS (S3, RDS, Glue, Lambda, EMR, Athena, EC2, IAM, CloudWatch), GCP (BigQuery, GCS, Cloud Composer, AlloyDB, PubSub)

BI & Visualization: Power BI (DAX), Tableau, Excel

Version Control & CI/CD: Docker, Kubernetes, Git, GitHub Actions, Terraform, Bitbucket, Linux

Frameworks: Streamlit

EDUCATION

Northeastern University, *Boston*

Sep 2024 - Dec 2025

Master's, Data Analytics (GPA: 3.85/4.0)

Coursework: Statistics, Intermediate Analytics, Big Data Management, Data Mining, Applications of Artificial Intelligence

PROJECTS

Weather Data Ingestion Pipeline (*Airflow, Docker, PostgreSQL*)

- Built an end-to-end data ingestion pipeline using Apache Airflow for orchestration, Docker for containerization, and **PostgreSQL** for storage, demonstrating cloud-native ETL design and database administration patterns.

PeakCart GCP Data Platform (*BigQuery, dbt, GCS, Cloud Composer, AlloyDB, Looker Studio*)

- Built Designed and built a production-grade Medallion Architecture (Bronze/Silver/Gold) on Google BigQuery using dbt Core, implementing SCD Type 2 dimensions, incremental fact tables, and a star schema covering 6 business entities, 92 automated data quality tests, and 5 advanced analytical SQL patterns including RFM segmentation, cohort analysis, and MERGE-based upserts.